

Cigarette Correlation

```
### Introduction
```

```
# Eleazar Lopez EID:el28789
```

```
# The datasets were two cigarette datasets with some variables including states, income,
```

```
# packs and cpi. I was interested in these datasets to see if there are any correlations in
```

```
# variables. I expect income and packs to have an inverse correlation.
```

```
# I pulled the datasets from github
```

```
# https://vincentarelbundock.github.io/Rdatasets/datasets.html
```

```
(https://vincentarelbundock.github.io/Rdatasets/datasets.html)
```

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com> (<http://rmarkdown.rstudio.com>).


When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
summary(cars)
```

```
##      speed      dist
##  Min.   : 4.0    Min.   : 2.00
##  1st Qu.:12.0    1st Qu.: 26.00
##  Median :15.0    Median : 36.00
##  Mean   :15.4    Mean   : 42.98
##  3rd Qu.:19.0    3rd Qu.: 56.00
##  Max.   :25.0    Max.   :120.00
```

Including Plots

You can also embed plots, for example:



Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

Tidy

```
#used library to load the tidyverse package
# datasets were already tidy
# used dplyr functions of rename to rename the x1 column to state
# used dplyr functions of select to get rid of the column of X1 which was useless
#used left_join and joined tidycigA and tidycigB by the variable of state and
#created the joinedcig dataset
library(tidyverse)
```

```
## — Attaching packages ————— tidyverse 1.3.1 —
```

```
## ✓ ggplot2 3.3.3      ✓ purrr   0.3.4
## ✓ tibble  3.1.6      ✓ dplyr   1.0.7
## ✓ tidyr   1.2.0      ✓ stringr 1.4.0
## ✓ readr   1.4.0      ✓ forcats 0.5.1
```

```
## — Conflicts ————— tidyverse_conflicts() —  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag() masks stats::lag()
```

```
library(readr)  
CigarettesB <- read_csv("CigarettesB.csv")
```

```
## Warning: Missing column names filled in: 'X1' [1]
```

```
##  
## — Column specification —————  
## cols(  
##   X1 = col_character(),  
##   packs = col_double(),  
##   price = col_double(),  
##   income = col_double()  
## )
```

```
CigarettesSW <- read_csv("CigarettesSW.csv")
```

```
## Warning: Missing column names filled in: 'X1' [1]
```

```
##  
## — Column specification —————  
## cols(  
##   X1 = col_double(),  
##   state = col_character(),  
##   year = col_double(),  
##   cpi = col_double(),  
##   population = col_double(),  
##   packs = col_double(),  
##   income = col_double(),  
##   tax = col_double(),  
##   price = col_double(),  
##   taxes = col_double()  
## )
```

```
tidycigB <- CigarettesB  
tidycigA <- CigarettesSW  
tidycigB <- tidycigB %>%  
  rename(state = X1)  
tidycigA <- tidycigA %>%  
  select(-c(X1))  
head(tidycigA)
```

```
## # A tibble: 6 × 9
##   state year  cpi population packs  income  tax price  taxes
##   <chr> <dbl> <dbl>      <dbl> <dbl>      <dbl> <dbl> <dbl> <dbl>
## 1 AL    1985  1.08   3973000  116.   46014968  32.5 102.   33.3
## 2 AR    1985  1.08   2327000  129.   26210736  37   101.   37
## 3 AZ    1985  1.08   3184000  105.   43956936  31   109.   36.2
## 4 CA    1985  1.08  26444000  100.  447102816  26   108.   32.1
## 5 CO    1985  1.08   3209000  113.   49466672  31    94.3  31
## 6 CT    1985  1.08   3201000  109.   60063368  42   128.   51.5
```

```
joinedcig <- left_join(tidycigA, tidycigB, by = "state")
head(joinedcig)
```

```
## # A tibble: 6 × 12
##   state year  cpi population packs.x  income.x  tax price.x  taxes packs.y
##   <chr> <dbl> <dbl>      <dbl> <dbl>      <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 AL    1985  1.08   3973000  116.   46014968  32.5  102.   33.3  4.96
## 2 AR    1985  1.08   2327000  129.   26210736  37    101.   37    5.11
## 3 AZ    1985  1.08   3184000  105.   43956936  31    109.   36.2  4.66
## 4 CA    1985  1.08  26444000  100.  447102816  26    108.   32.1  4.50
## 5 CO    1985  1.08   3209000  113.   49466672  31     94.3  31    NA
## 6 CT    1985  1.08   3201000  109.   60063368  42    128.   51.5  4.67
## # ... with 2 more variables: price.y <dbl>, income.y <dbl>
```

Exploratory Data Analysis

```
# used select function to get rid of the variable state and year as they were not
# necessary
# used function cor to see the correlation between all variables
# used chart.Correlation to view correlation histograms and correlation coefficients
# for all variables
joinedcig <- joinedcig %>%
  select(-c(year, state))
cor(joinedcig)
```

```
##               cpi population   packs.x   income.x       tax   price.x
## cpi           1.00000000 0.04758017 -0.4994643 0.2317893 0.6857145 0.9116556
## population    0.04758017 1.00000000 -0.2112834 0.9573113 0.1659856 0.1458604
## packs.x       -0.49946432 -0.21128337 1.0000000 -0.3317847 -0.6421176 -0.6524732
## income.x      0.23178932 0.95731126 -0.3317847 1.0000000 0.3372751 0.3375339
## tax           0.68571446 0.16598557 -0.6421176 0.3372751 1.0000000 0.8993727
## price.x       0.91165558 0.14586043 -0.6524732 0.3375339 0.8993727 1.0000000
## taxes         0.70412144 0.18891721 -0.6574167 0.3582307 0.9853330 0.9203278
## packs.y               NA               NA               NA               NA               NA
## price.y              NA               NA               NA               NA               NA
## income.y             NA               NA               NA               NA               NA
##               taxes packs.y price.y income.y
## cpi           0.7041214         NA         NA         NA
## population    0.1889172         NA         NA         NA
## packs.x       -0.6574167         NA         NA         NA
## income.x      0.3582307         NA         NA         NA
## tax           0.9853330         NA         NA         NA
## price.x       0.9203278         NA         NA         NA
## taxes         1.0000000         NA         NA         NA
## packs.y               NA          1         NA         NA
## price.y              NA         NA          1         NA
## income.y             NA         NA         NA          1
```

```
library(PerformanceAnalytics)
```

```
## Loading required package: xts
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
##
## Attaching package: 'xts'
```

```
## The following objects are masked from 'package:dplyr':
##
##   first, last
```

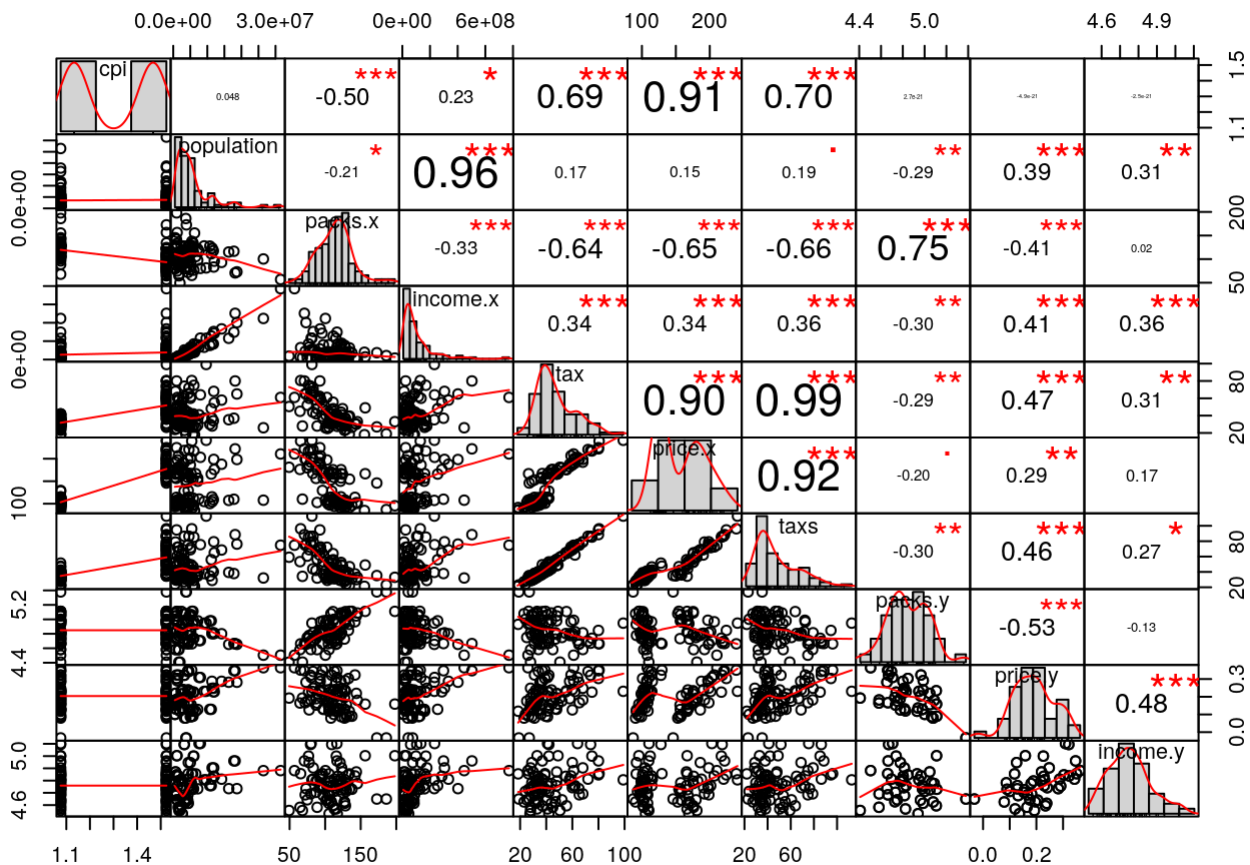
```
##
## Attaching package: 'PerformanceAnalytics'
```

```
## The following object is masked from 'package:graphics':
```

```
##
```

```
## legend
```

```
chart.Correlation(joinedcig, histogram = TRUE, method = "pearson")
```




The most correlated variables are tax and taxes but if we ignore that due to them both being a tax variable the next most correlated variables is income.x and population. The least correlated variables are packs.x and taxes. As packs increase the income tends to decrease.

Clustering

```
# used fviz_nbclust with the joinedcig dataframe to see the optimal number of clusters
# used function pam and named it pam_results
# used ggpairs to visualize the clusters
# with this vizualization we can see which variables have the strongest correlation,
# positive and negative
# we can more clearly see how tax and taxes have the strongest correlation between all
# variables
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(cluster)
fixedjoinedcig <- joinedcig %>%
  select(-c(packs.y, price.y, income.y))
fviz_nbclust(fixedjoinedcig, pam, method = "silhouette")
```



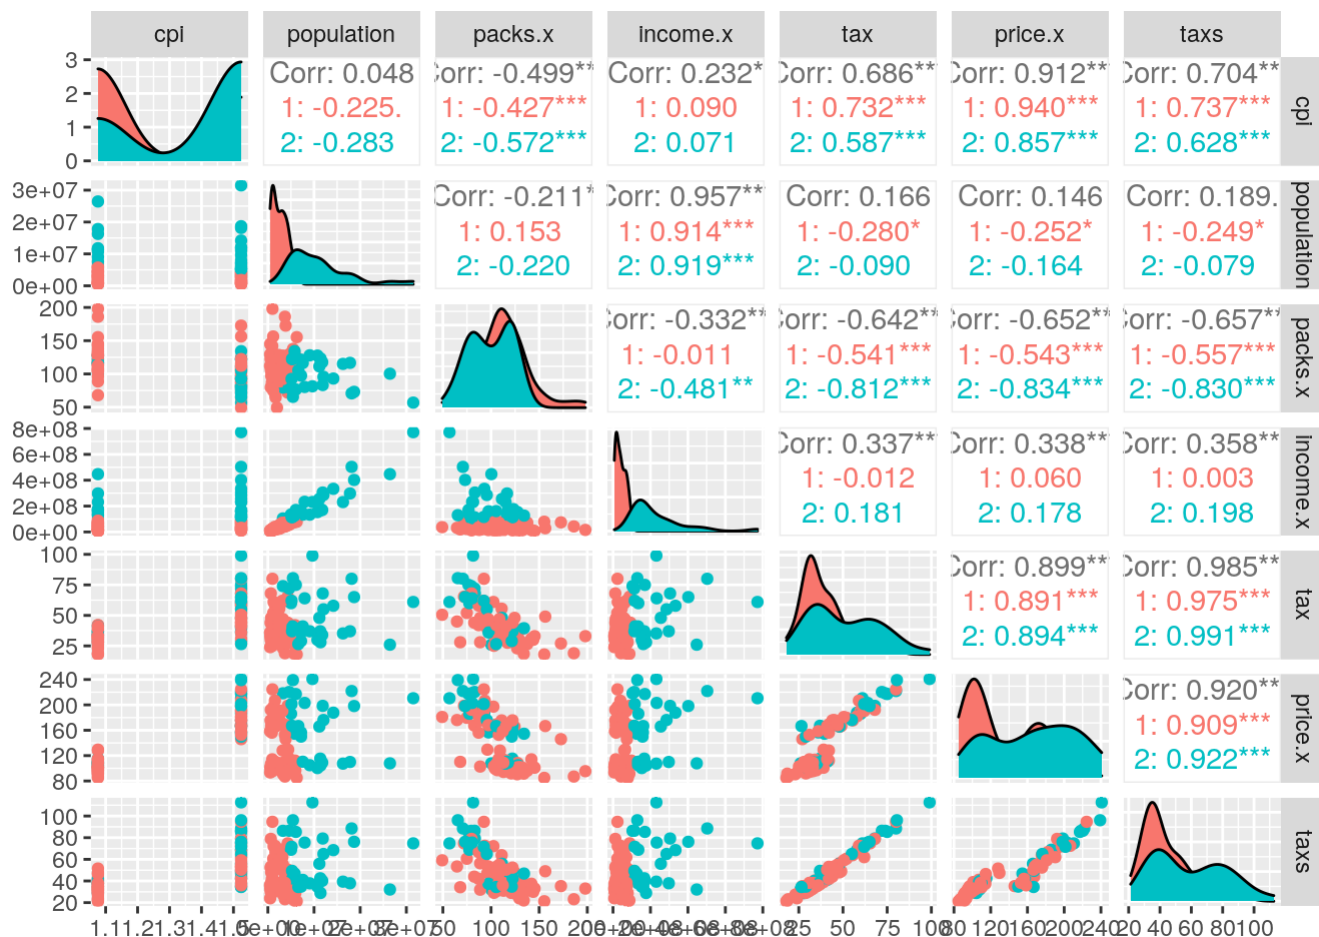
```
pam_results <- pam(fixedjoinedcig,2)
pam_results
```

```
## Medoids:
##      ID  cpi population  packs.x  income.x tax price.x  taxes
## [1,] 10 1.076    2830000 113.7456  37902896  34 101.842  37.917
## [2,]  8 1.076    11352000 122.1811 166919248  37 115.290  42.490
## Clustering vector:
## [1] 1 1 1 2 1 1 1 2 1 1 1 2 1 1 1 1 1 2 1 1 1 1 1 1 1 2 1 1 2 2 1 1 2 1 1
## [39] 1 1 2 1 1 1 1 1 1 1 1 1 1 2 1 2 1 2 2 1 1 2 2 1 1 1 2 2 1 2 2 2 1 1 2 1 1 1
## [77] 2 1 1 2 2 1 1 2 1 1 1 2 2 1 2 1 2 2 1 1
## Objective function:
##      build      swap
## 46995509 44449436
##
## Available components:
## [1] "medoids"      "id.med"      "clustering"  "objective"   "isolation"
## [6] "clusinfo"     "silinfo"     "diss"        "call"        "data"
```

```
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##      method from
##      +.gg      ggplot2
```

```
fixedjoinedcig %>% mutate(cluster = as.factor(pam_results$clustering)) %>%
  ggpairs(columns = 1:7, aes(color = cluster))
```

Dimensionality Reduction

```
# used dyplr function of select to get rid of packs.y, price.y, and income.y
# used function prcomp to find principal components
# used the function cut to change cpi to a categorical variable
# used fviz_cluster to separate the clusters by cpi
fixedjoinedcig <- joinedcig %>%
  select(-c(packs.y, price.y, income.y))
prcomp(fixedjoinedcig)
```

```
## Standard deviations (1, .., p=7):
## [1] 1.206537e+08 1.571692e+06 3.818919e+01 1.792198e+01 8.807293e+00
## [6] 1.993977e+00 4.766130e-02
##
## Rotation (n x k) = (7 x 7):
##
```

	PC1	PC2	PC3	PC4
cpi	4.319543e-10	-8.641615e-08	-3.333794e-03	-2.831933e-03
population	4.318892e-02	9.990669e-01	-1.993289e-05	-2.220790e-06
packs.x	-7.109888e-08	6.058607e-06	4.028361e-01	-9.137535e-01
income.x	9.990669e-01	-4.318892e-02	1.021067e-06	8.828357e-08
tax	4.507149e-08	-5.575065e-06	-2.864985e-01	-9.478971e-02
price.x	1.226514e-07	-1.712943e-05	-7.938283e-01	-3.774543e-01
taxs	5.735084e-08	-6.556501e-06	-3.542134e-01	-1.165750e-01

```
##
```


	PC5	PC6	PC7
cpi	-1.023253e-02	-9.556602e-03	-9.998924e-01
population	-9.157836e-07	-2.888839e-07	-1.444217e-09
packs.x	5.247904e-02	4.813206e-03	6.617956e-04
income.x	3.791183e-08	1.132105e-08	1.360896e-11
tax	6.154963e-01	-7.280757e-01	1.883610e-03
price.x	-4.750573e-01	-4.006410e-02	8.960266e-03
taxs	6.265989e-01	6.842413e-01	-1.144094e-02

```
fixedjoinedcig$cpi <- cut(fixedjoinedcig$cpi,
                          breaks=c(1, 1.5),
                          labels=c("A"))
fviz_cluster(pam_results, data = fixedjoinedcig,
              shape = fixedjoinedcig$cpi) +
  geom_point(aes(shape = fixedjoinedcig$cpi)) +
  guides(shape = guide_legend(title = "shape"))
```

```
## Warning in if (shape %in% colnames(data)) {: the condition has length > 1 and
## only the first element will be used
```

```
## Warning: Removed 48 rows containing missing values (geom_point).
```

```
## Warning: Removed 48 rows containing missing values (geom_point).
```



Classification & Cross-Validation

```
# made new data set called ffjoinedcig from fixedjoinedcig by mutating and creating a new
# variable Economic_status rich was over 100,000,000 in income
# made a table with the actual and prediction of ffjoinedcig
# created an ROC plot that showed that economic_status is a perfect indicator of income
library(tidyverse)
ffjoinedcig <- fixedjoinedcig %>%
  mutate(Economic_status = ifelse(income.x > 100000000, "rich", "not rich"))
actual <- ffjoinedcig$income.x
prediction <- ffjoinedcig$Economic_status
table(actual = ffjoinedcig$income.x, prediction = ffjoinedcig$Economic_status) %>%
  addmargins()
```

##		prediction		
##	actual	not	rich	Sum
##	6887097	1	0	1
##	7116756	1	0	1
##	8340000	1	0	1
##	8672948	1	0	1
##	9785230	1	0	1
##	9927301	1	0	1
##	10293195	1	0	1
##	11577261	1	0	1
##	12243384	1	0	1
##	12448607	1	0	1
##	14229156	1	0	1
##	14454129	1	0	1
##	14575292	1	0	1
##	14581495	1	0	1
##	15767469	1	0	1
##	16296835	1	0	1
##	17258916	1	0	1
##	18237436	1	0	1
##	19462380	1	0	1
##	20852964	1	0	1
##	21778072	1	0	1
##	22868920	1	0	1
##	23786644	1	0	1
##	25045934	1	0	1
##	25678534	1	0	1
##	26210736	1	0	1
##	28649564	1	0	1
##	31716160	1	0	1
##	32611268	1	0	1
##	34784360	1	0	1
##	36205164	1	0	1
##	36293064	1	0	1
##	37278220	1	0	1
##	37902896	1	0	1
##	38536176	1	0	1
##	39377292	1	0	1
##	42703144	1	0	1
##	43395580	1	0	1
##	43956936	1	0	1
##	45995496	1	0	1
##	46014968	1	0	1
##	46241956	1	0	1
##	49466672	1	0	1
##	53431900	1	0	1
##	56626672	1	0	1
##	57749668	1	0	1
##	60063368	1	0	1
##	60170928	1	0	1
##	63152360	1	0	1
##	63333300	1	0	1

##	64846548	1	0	1
##	65732720	1	0	1
##	69341920	1	0	1
##	71209312	1	0	1
##	71751616	1	0	1
##	72050072	1	0	1
##	74079712	1	0	1
##	74851664	1	0	1
##	78364336	1	0	1
##	79104656	1	0	1
##	83903280	1	0	1
##	84572688	1	0	1
##	87361632	1	0	1
##	88870496	1	0	1
##	92946544	1	0	1
##	98328688	1	0	1
##	104315120	0	1	1
##	113216856	0	1	1
##	114259984	0	1	1
##	115959680	0	1	1
##	117639672	0	1	1
##	126525008	0	1	1
##	129680832	0	1	1
##	133549208	0	1	1
##	133728040	0	1	1
##	135115456	0	1	1
##	153455776	0	1	1
##	157633568	0	1	1
##	159800448	0	1	1
##	161441792	0	1	1
##	166919248	0	1	1
##	170033840	0	1	1
##	170051568	0	1	1
##	176786352	0	1	1
##	231003152	0	1	1
##	231594240	0	1	1
##	233208576	0	1	1
##	255312928	0	1	1
##	285923232	0	1	1
##	297728512	0	1	1
##	304767456	0	1	1
##	333525344	0	1	1
##	402096768	0	1	1
##	447102816	0	1	1
##	503163328	0	1	1
##	771470144	0	1	1
##	Sum	66	30	96

```
F1 <- function(y_hat, y, positive){
  sensitivity <- mean(y_hat[y == positive] == positive)
  precision <- mean(y[y_hat == positive] == positive)
  2*(sensitivity*precision)/(sensitivity + precision)
}

F1(prediction, actual, "rich")
```

```
## [1] NaN
```

```
library(plotROC)
ROC <- ggplot(ffjoinedcig) +
  geom_roc(aes(d = Economic_status, m = income.x))
ROC
```

```
## Warning in verify_d(data$d): D not labeled 0/1, assuming not rich = 0 and rich =
## 1!
```

